Smoothed complexity of local max-cut for special graphs

November 30, 2017

Abstract

The problem of finding a max-cut in graphs has many applications and motivation. As this problem to be NP-hard, research was made on simpler problems such as a local max-cut.

1 Introduction

Let $G = (V, E)$ be a graph with n vertices and m edges. Let $w : E \to [-1, 1]$ be an edge weight function. The local max-cut problem consists in finding a partition of the vertices σ such that the total cut weight, defined as :

$$h(\sigma) = \frac{1}{2} \sum_{uv \in E} (1 - \sigma(u)\sigma(v))w(uv)$$

is locally optimal. By locally optimal, we mean that there exists no vertex v such that, if we flip the vertex, i.e. change the sign of $\sigma(v)$, then $h(\sigma)$ increases.

A really naive algorithm called FLIP solves this problem. It finds a vertex which if flipped leads to an amelioration of the total weight cut, flips it and repeat until no such vertex exists. Some instances of graphs have been found to have an exponential number of steps before terminating. However, for most of the graphs, FLIP terminates in a reasonable time. Moreover, when adding a small amount of noise to the complicated graphs, FLIP's running time improved greatly. This lead to the study of the smoothed complexity of FLIP, in which noise is added to the edge weights.

Etscheid and Röglin (2014) [**?**] proved that this complexity was at most quasi-polynomial in n for arbitrary graphs, with the insight that it may be polynomial. Angel et al. (2016) [**?**] proved that the complexity was polynomial for complete graphs.
We study here other special graphs, for which the complexity is polynomial with the hope that it would extend for general graphs.

2 Notation and preliminary lemmas

Let $X = (X_e)_{e \in E} \in [-1,1]^E$ a random vector with independent entries, corresponding to the edge weights. We assume that X_e has density f_e with respect to the Lebesgue mesasure, and we denote $\phi = max_{e \in E} ||f_e||_\infty$.

Lemma 2.1 *(Lemma 2.1 [?])*
Let $\alpha_1, ..., \alpha_k$ be k linearly independent vectors in $\mathbb{Z}^E$. Then the joint density of$(\langle \alpha_i, X \rangle)_{i \le k}$ is bounded by ϕ^k. In particular, if sets $J_i \in \mathbb{R}$ have measure at most ϵ each, then

$$\mathcal{P}(\forall i \in [k], \langle \alpha_i, X \rangle \in J_i) \le (\phi \epsilon)^k$$

We define a move vector α_v as a vector indexed by E whose entries are :

$$\alpha_{uw} = \begin{cases} \sigma(u)\sigma(w) & \text{if } uw \in E \text{ and } ((u = v) \text{ or } (w = v)) \\ 0 & \text{otherwise} \end{cases}$$

For a sequence $L = (v_1, ..., v_l)$ of l moves and initial state σ_0, let $\alpha_i, i \in [l]$ be the corresponding move vectors. Let σ_t be the state just after flip of vertex v_t. We define matrix A_L as the concatenation of the move vectors as columns. We call a sequence ϵ-slowly improving if all moves yield an improvement of at most ϵ.

We introduce here the concept of critical block. A block B is defined as a substring of a sequence L.

Let $S(L)$ be the set containing the distinct vertices in L and $s(L)$ be its cardinality, $s_1(L)$ the number of distinct vertices in L that appear only once, $s_2(L)$ the number of distinct vertices in L that appear multiple times. Let $l(B)$ be the length of the block, i.e. the number of moves. A block B is critical if $l(B) \ge (1+\beta)s(B)$ and every block B' strictly contained in B has $l(B') < (1+\beta)s(B')$.

Lemma 2.2 *(Lemma 4.1 [?])*
For complete graphs with n vertices: fix any positive integer $n \ge 2$ and a constant $\beta > 0$. Given a sequence consisting of $s(L) < n$ letters and with length $l(L) \ge (1+\beta)s$, there exists a critical block B in L. Moreover, a critical block satisfies $l(B) = \lceil (1+\beta)s(B) \rceil$. Moreover $rank(B) \ge \dfrac{1+4\beta}{1+3\beta} s(B)$

Lemma 2.3
If $s(L) < n$, and $rank(L) \geq (1+\theta)s(L)$ then

$$P(L \text{ is } \epsilon\text{-slowly improving from some sigma}_0) \leq (2n/\epsilon)^{s_0}(64\phi\epsilon)^{rank(A_L)}$$

Where s_0 is the number of vertices that have at least one neighbor which is not in L.

Proof. Let I be the set of the edges corresponding to independent rows in A_L. They do not depend on σ_0 since the starting configuration only multiplies each row by 1 or -1.
Define $T = \{v \in V : vw \in I \text{ or } v \in S(L)\}$. $|T| \leq 2rank(A_L) + s(L)$.
We split $h(\sigma)$ in three part $h_0(\sigma), h_1(\sigma), h_2(\sigma)$ Where h_i is the restriction of h to the edges that have i endpoints in T.
$h_0(\sigma_t) - h_0(\sigma_{t-1}) = 0$. Since the edges whose both endpoints are not flipped do not provoke a change in the total weight cut.

$$h_1(\sigma_t) - h_1(\sigma_{t-1}) = -\sigma_t(v_t) \sum_{u \notin T, uv_t \in E} \sigma_0(u) X_{uv_t}$$
$$= \sigma_t(v_t) Q(v_t).$$

where $(Q(v_t) = -\sum_{u \notin T, uv_t \in E} \sigma_0(u) X_{uv_t})$. Since $|X_e| \leq n$ and the maximum neighbours of a vertex is n, $Q(v_t) \in [-n, n]$. By defining $D = 2\epsilon\mathbb{Z} \cap [-n, n]$, there exists some $d(v_t) \in D$ such that $|Q(v_t) - d(v_t)| \leq \epsilon$.

$$h_2(\sigma_t) - h_2(\sigma_{t-1}) = \langle \alpha', X \rangle \text{ where}$$
$$\alpha'_t = \begin{cases} -\sigma_t(u)\sigma_t(w) & uw \in E, v_t \in \{u, w\}, \{u, w\} \subseteq T \\ 0 & \text{otherwise} \end{cases}$$

Since α'_t concerns only the rows linearly independant from A_L, $rank([\alpha'_t]_{t \leq l}) = rank(A_L)$.

$$h(\sigma_t) - h(\sigma_{t-1} = \langle \alpha', X \rangle + \sigma_t(v_t)d_t + \delta_t \text{ where } |\delta_t| \leq \epsilon$$

Since $|\delta_t| \leq \epsilon$, L is ϵ-slowly improving implies that:

$$|\langle \alpha', X \rangle + \sigma_t(v_t)d_t| \leq 2\epsilon \quad \forall t \leq l$$

We need that $\langle \alpha', X \rangle \in [-d_t - 2\epsilon, -d_t + 2\epsilon] \cup [d_t - 2\epsilon, d_t + 2\epsilon]$, Using lemma 2.1, this is at most $8\phi\epsilon^{rank(A_L)}$. Using union bound over $\sigma_{t \in T}$ and d :

$$\text{P(L is } \epsilon\text{-slowly improving from some } \sigma_0 \leq 2^{2rank(A_L)+s}(\frac{2n}{\epsilon})^{s_0}(8\phi\epsilon)^{rank(A_L)}$$

Remark the s_0 instead of s in exponent, since vertices that have no-non flipped neighbors need not be taken in this bound.
By using $rank(L) \geq (1+\theta)s(L)$, we get the desired bound.

3 Proof for graphs with one clique and low degrees vertices

Let $G = (V, E)$ be a graph with n vertices and m edges. Assume that this graph contains a clique H of r vertices and that the degree of vertices in the set $G \setminus H$ is at most $log(n)$.

Proposition 3.1 *With high probability, there exists no ϵ-slowly improving sequence of length $2n$ from any starting configuration σ_0, for ϵ is O(1/poly(n)).*

Proving this proposition implies that the smoothed complexity is O(poly(n)). If there exists no such sequence, then $2n^2/\epsilon$ sequence of $2n$ moves yield an improvement of at least $2n^2$ which is the maximum improvement possible since $h(\sigma) \in [-n^2, n^2]$.
Thus number of steps is $O(n^3/\epsilon)$ which is $O(poly(n))$.

We will prove the proposition by considering different sequences of size 2n.
Let $p > 1$, l(L) = 2n :
E is the event corresponding to $\exists L, \sigma_0,$ s.t. L is ϵ-slowly improving from σ_0
E_1 is the event corresponding to $\exists L, \sigma_0,$ s.t. L is ϵ-slowly improving from σ_0 and $S(L) \nsubseteq H$.
E_2 is the event corresponding to $\exists L, \sigma_0,$ s.t. L is ϵ-slowly improving from σ_0 and $S(L) \subseteq H$ and $s(L) < r$.
E_3 is the event corresponding to $\exists$ critical block $B, \sigma_0,$ s.t. B is ϵ-slowly improving from σ_0 and $S(L) \subseteq H$, $s(B) = r$ By the lemma 2.2 on existence of critical blocks and the fact that if $s(L) = n$ implies that some vertex with degree at most log(n) is chosen we can have this bound:

$$P(E) \leq P(E_1) + P(E_2) + P(E_3)$$

Consider E_1. Fix L and σ_0 then there must be a vertex $v \in S(L)$ whose degree is at most $log(n)$. By lemma 2.1, the probability that $\alpha_v \in [0, \epsilon] = \phi\epsilon$. Using union bound on the number of vertices and the starting configuration of those vertices we have:

$$P(E_1) \leq 2^{log(n)} n\phi\epsilon \tag{1}$$

By taking$\epsilon = n^{-2-\eta}\phi^{-1}$ with $\eta > 0$ there exists no such sequence with high probability.

Now consider E_2. We consider the subgraph G' induced by the restriction to the clique H. We observe that $rank'_G(A_L) \leq rank_G(A_L)$ Since $G'(A_L)$ is a submatrix of $G(A_L)$.
By lemma 2.2, there exists a critical block B whose rank is at least 1.25 s(B).We now can use lemma **??** to have this bound:

$$P(B \text{ is } \epsilon\text{-slowly improving from some } \sigma) \leq 2(\frac{4n}{\epsilon})^{s(B)}(8\phi\epsilon)^{\frac{5s(B)}{4}}$$

Because the number of blocks using s letters is n^{2s}, we have:

$$P(E_2) \leq 2\sum_{s<n}(64\phi^{5/4}n^3\epsilon^{1/4})^s$$

By choosing $\epsilon = n^{-(12+\eta)}\phi^{-5}$ This sums goes to zero.

For E_3, we use a trick to show that the $rank(A_L) \geq 1.25s(B) - (s(B) - s_0(B))$. We choose some $w \in V \setminus H$. For each vertex v which has a non-flipped neighbour, we delete that edge and add vw to the graph. This does not change $rank(A_L)$ since the row added is the same as the row deleted times 1 or -1. Now we add edges from w to the remaining vertices on the clique, increasing the rank by at most $s(B) - s_0(B)$. The subgraph G' determined by $H \cup \{w\}$ is thus complete and we can use lemma 2.2 to have $rank_{G'}(A_L) \geq 1.25s(B)$. Then $rank_G(A_L) \geq 1.25s(B) - (s(B) - s_0(B))$. By lemma 2.3 we have :

$$P(E_4) \leq \sum_{s<n} n^{2s}(\frac{2n}{\epsilon})^{s_0}(64\phi\epsilon)^{s/4+s_0} \leq \sum_{s<n}(Cn^3\phi^{5/4}\epsilon^{1/4})^s$$

By choosing $\epsilon = n^{-(12+\eta)}\phi^{-5}$ This sums goes to zero.